

# 1 AccessRefinery: Fast Mining Concise Access Control Intents 2 on Public Cloud

3 NING KANG, Xi'an Jiaotong University, China

4 PENG ZHANG\*, Xi'an Jiaotong University, China

5 JIANYUAN ZHANG, Xi'an Jiaotong University, China

6 HAO LI, Xi'an Jiaotong University, China

7 DAN WANG, Xi'an Jiaotong University, China

8 ZHENRONG GU, Xi'an Jiaotong University, China

9 WEIBO LIN, Huawei Cloud, China

10 SHIBIAO JIANG, Huawei Cloud, China

11 ZHU HE, Huawei Cloud, China

12 XU DU, Huawei Cloud, China

13 LONGFEI CHEN, Huawei Cloud, China

14 JUN LI, Huawei Cloud, China

15 XIAOHONG GUAN, Xi'an Jiaotong University, China

16 Modern cloud applications heavily rely on Identity and Access Management (IAM) services to enforce flexible  
17 access control over their data. However, the flexibility comes at a cost: *IAM policies* are often complex and  
18 prone to misconfigurations, leading to risks of data exposure. There is an increasing need to mine a compact  
19 set of intents that describe what the policies collectively try to achieve, thereby enabling operators to better  
20 understand their policies. However, existing tools on mining access control intent have two limitations: (1)  
21 the mining process is *slow* and even times out on some complex policies; (2) the mined intents are *excessive* in  
22 number and thus still hard to understand. To overcome these, this paper presents *AccessRefinery*, which can  
23 speed up the mining process while reducing the number of intents. The key idea for the speedup is to reduce  
24 the redundancy of the multi-round SMT solving, by preprocessing the constraints into bit-vector constraints.  
25 For intent reduction, *AccessRefinery* computes a compact set of intents that can cover the mined intents, by  
26 solving a *min-set-cover* problem. Experiments based on real and synthetic datasets show that *AccessRefinery*  
27 achieves a  $\sim 10\text{--}100\times$  speedup in intent mining, and reduces the number of intents by up to  $\sim 10\times$ .  
28

29  
30 CCS Concepts: • Security and privacy → Access control; Logic and verification.

31 Additional Key Words and Phrases: Cloud computing, access control, intent mining, SMT

---

32 \*Corresponding author.

33  
34 Authors' Contact Information: Ning Kang, Xi'an Jiaotong University, Xi'an, Shaanxi, China, kangning@stu.xjtu.edu.cn;  
35 Peng Zhang, Xi'an Jiaotong University, Xi'an, Shaanxi, China, p-zhang@xjtu.edu.cn; Jianyuan Zhang, Xi'an Jiaotong  
36 University, Xi'an, Shaanxi, China, 2244313675@stu.xjtu.edu.cn; Hao Li, Xi'an Jiaotong University, Xi'an, Shaanxi, China,  
37 hao.li@xjtu.edu.cn; Dan Wang, Xi'an Jiaotong University, Xi'an, Shaanxi, China, dan-wang@stu.xjtu.edu.cn; Zhenrong Gu,  
38 Xi'an Jiaotong University, Xi'an, Shaanxi, China, zrgu@stu.xjtu.edu.cn; Weibo Lin, Huawei Cloud, Xi'an, Shaanxi, China,  
39 linweibo@huawei.com; Shibiao Jiang, Huawei Cloud, Xi'an, Shaanxi, China, jiangshibiao1@huawei.com; Zhu He, Huawei  
40 Cloud, Xi'an, Shaanxi, China, hezhu4@huawei.com; Xu Du, Huawei Cloud, Xi'an, Shaanxi, China, duxu1@huawei.com;  
41 Longfei Chen, Huawei Cloud, Xi'an, Shaanxi, China, chenlongfei@huawei.com; Jun Li, Huawei Cloud, Xi'an, Shaanxi, China,  
42 China, alanlee@huawei.com; Xiaohong Guan, Xi'an Jiaotong University, Xi'an, Shaanxi, China, xhguan@xjtu.edu.cn.

43  
44 Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee  
45 provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and  
46 the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be  
47 honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists,  
48 requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

49 © 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM XXXX-XXXX/2026/2-ART

<https://doi.org/XXXXXXXX.XXXXXXXX>

50      **ACM Reference Format:**

51      Ning Kang, Peng Zhang, Jianyuan Zhang, Hao Li, Dan Wang, Zhenrong Gu, Weibo Lin, Shibiao Jiang, Zhu  
 52      He, Xu Du, Longfei Chen, Jun Li, and Xiaohong Guan. 2026. AccessRefinery: Fast Mining Concise Access  
 53      Control Intents on Public Cloud. 1, 1 (February 2026), 23 pages. <https://doi.org/XXXXXX.XXXXXXX>

54      **1 Introduction**

55      Software applications are increasingly deployed in the cloud [4, 5, 16, 17, 31, 41]. To secure these  
 56      applications, major cloud providers, including Amazon Web Services (AWS) [5], Microsoft Azure  
 57      [31], and Google Cloud Platform (GCP)[25] offer Identity and Access Management (IAM) systems.  
 58      These systems follow a shared responsibility model: customers secure their services by writing  
 59      *access control policies* (or *IAM policies*) while the cloud provider's engine evaluates requests against  
 60      policies to determine whether access should be authorized [18, 41].

61      To enable fine-grained access control, *IAM policies* support flexible features including wildcard  
 62      matching (e.g., regular expressions) and Boolean logic (e.g., conjunctions, disjunctions, and nega-  
 63      tions) [4, 5, 16–18, 31, 41]. However, such flexibility also makes *IAM policies* complicated and error-  
 64      prone. Misconfigured *IAM policies* can accidentally allow unintended access, exposing millions of  
 65      customers' data to the public [37, 40, 44, 46].

66      A promising way to handle the complexity of *IAM policies* is to mine a concise summary of  
 67      *intents*, i.e., a compact set of declarative statements that describe *who* can perform *which* actions on  
 68      *what* sources (see §2.2). By reviewing such intents, users can more confidently determine whether  
 69      the policies correctly grant the intended access [4], thereby detecting potential misconfigurations.  
 70      Users can also better understand their policies and refine their designs [16]. Moreover, mining  
 71      *access control intents* is crucial to software engineering. First, many software systems today are  
 72      deployed to the cloud, where access is safeguarded by access control policies. Mining the intents of  
 73      these policies can formally validate that access is granted as intended. Second, it is closely related  
 74      to specification mining, which aims to infer invariants or specifications of given programs [19, 22].

75      The following shows the desired goals for intent mining (the first three are from [4]).

- 76      • **Soundness.** All requests allowed by the policy should be included in the intents.
- 77      • **Precision.** Any request not allowed by the policy should be excluded as much as possible.
- 78      • **Conciseness.** The number of intents should be minimal so that they are concise and easy  
     80      to understand.
- 81      • **Speed.** The mining process should finish within a reasonable time (e.g., on a level of sec-  
     83      onds), so as to be responsive to frequent user requests.

84      Among these goals, soundness is a hard requirement, since we do not want to overlook any (un-  
 85      intended) allowed access. Moreover, precision and conciseness are soft requirements and they may  
 86      conflict with each other. For example, consider two intents: (Action: {list, create}, Principal:  
 87      {user1}), and (Action: {list}, Principal: {user1, user2}). Merging them into ({list, create},  
 88      {user1, user2}) sacrifices precision, while keeping them separate sacrifices conciseness.

89      While the state-of-the-art mining tool *Access Analyzer* (commercially deployed) [4] proposed  
 90      by AWS meets the first two goals, it still has some limitations in the conciseness and speed.

91      **Limitation 1: Speed.** *Access Analyzer* encodes both *IAM policies* and each candidate intent (a  
 92      hypothesized intent to be checked) as a set of Satisfiability Modulo Theories (SMT) constraints,  
 93      which are solved using off-the-shelf SMT solvers (also widely applied to analyze *IAM policies* [4, 5,  
 94      7, 8, 16–18, 31, 41]). If the constraints are satisfiable, the candidate intent is confirmed as an intent.  
 95      Since each intent is a combination of values defined over multiple and possibly different domains  
 96      (e.g., strings, bit-vectors, integers), the SMT solving process needs to reason about a product space

of values of each domain, which can be quite large. Moreover, consider a policy with  $n$  keys and  $m$  distinct values for these keys. The number of candidate intents can be up to  $n \times m$ . Validation of each intent requires one round of SMT solving. To reduce the number of rounds for SMT solving, *Access Analyzer* uses a stratified approach: start from a most coarse-grained intent (all wildcards), and gradually divide it into fine-grained ones, if it can be covered by any of them (see §2.3 for a detailed discussion). Although this method avoids enumeration, the number of solving rounds is still large. As we will show in Table 2, the number of SMT solving can reach  $O(10^3)$  for a policy with 5 keys and 15 statements. Considering each round of SMT solving costs  $O(10)$  to  $O(1000)$ ms, the mining process can take tens of seconds or several minutes, and even times out after 1 hour (see Fig. 13).

**Limitation 2: Conciseness.** The number of intents returned by *Access Analyzer* can still be too large for users to understand. To make the intent concise, *Access Analyzer* has applied some simple reduction rules. For example, suppose there are two intents (*Action* : *list*, *IpAddress* : 112.0.0.0/24) and (*list*, 112.0.0.0/25), then (*list*, 112.0.0.0/25) is removed since it is a subset of (*list*, 112.0.0.0/24). Such a simple reduction rule does not fully account for the redundancy among the intents. Suppose there are three intents (*Action* : {*list*, *create*}, *IpAddress* : 112.0.0.0/25), (*list*, *IpAddress* : 112.0.0.0/24), and (*create*, *IpAddress* : 112.0.0.0/24). The first intent is jointly covered by the second and third intents, and therefore should be eliminated. As we will show, without considering such redundancy, we may end up with 100 intents for an *IAM policy* with just 10 statements, where the minimum number of intents is actually 10 (see Fig. 12).

To overcome the above two limitations, we propose *AccessRefinery*, a new intent mining tool for *IAM policies*, which is faster, and more importantly guarantees the conciseness of intents. Specifically, we make the following contributions:

**(1) We propose a method to speed up the mining process by preprocessing SMT constraints.** we observe that even though the mining process involves multi-round SMT solving, the possible values of variables are known in advance. We leverage this observation by partitioning each domain into a set of *equivalence classes* (ECs), each of which is represented as an integer, such that the combinations of values from multi-domain values can be represented as simple bit-vectors. As a result, the original SMT problem can be reduced to an SAT problem, which can be solved faster without invoking theory solvers, e.g., string solver, bit-vector solver, integer solver. Moreover, the multiple rounds of SMT solving share intermediate results that can be reused for speedup. Crucially, the preprocessing is per mining process: it needs to be performed once and can be shared by multiple rounds of SMT solving. We realize the above idea with the *Multi-Theory Constraint Preprocessor* (MCP), which can uniformly handle multiple domain types of *IAM policies* including regular expressions, prefixes, ranges, and enumerations.

**(2) We propose a method to efficiently reduce the number of intents by solving a min-set-cover problem.** Recall that the soundness goal requires that the space of all requests specified by intents should cover the space of requests allowed by the policy. Therefore, computing the minimum number of intents can be formulated as a *min-set-cover* problem [29], i.e., finding the minimum number of intents that can cover the space of requests allowed by the policy. Note here, the space of requests is actually a product space for multiple domains, where a domain can be infinite, e.g., Resource: “dept\*/user1.txt”. Thanks to the preprocessing, we have reduced each domain into a finite set of ECs. Here, each EC represents a group of requests that exhibit the same behavior under both the policy and the intents; that is, requests within an EC are either all matched by the same set of intents and policy, or none are. We partition the space of the policy and intents into ECs (unlike MCP, this partitioning occurs across multiple domains), enabling both

148 the policy and intents to be represented as finite sets of ECs. With this representation, the *min-set-*  
 149 *cover* problem is transformed into one over a finite set; that is, the objective becomes selecting the  
 150 smallest number of intents whose union covers all ECs allowed by the policy. However, computing  
 151 these ECs requires iteratively solving SMT problems, which is quite slow (see §3.1). Inspired by [49,  
 152 51], we choose to use *Binary Decision Diagrams* (BDDs) [2], which allow incremental computations  
 153 of ECs.

154 **(3) We implemented a new mining  
 155 tool named *AccessRefinery*, and used  
 156 both real and synthetic datasets to  
 157 show its effectiveness (Both the tool  
 158 and the synthetic dataset are open-  
 159 sourced).** Compared with existing meth-  
 160 ods based on SMT solvers, our method  
 161 speeds up the mining process by  $\sim 10\times$  to  
 162  $100\times$  on synthetic policies, and by  $\sim 10\times$   
 163 on real policies. Moreover, our method re-  
 164 duces the number of intents by  $\sim 10\times$ . Fur-  
 165 thermore, we design MCP as a data struc-  
 166 ture that we believe will also benefit other  
 167 studies.

## 168 2 Motivation

### 169 2.1 Background on IAM Policy

170 *IAM policies* are typically *attribute-based*  
 171 (e.g., AWS) or *role-based* (e.g.,  
 172 Azure), differing in syntax but semanti-  
 173 cally equivalent [18]. This paper focuses  
 174 on the AWS attribute-based paradigm [5].

175 Fig. 2 shows the simple syntax for the  
 176 AWS policy language as defined in [5, 8].  
 177 A policy is a list of statements. Each state-  
 178 ment is a quintuple: (*Effect*, *Principal*,  
 179 *Action*, *Resource*, *Condition*). A request is  
 180 allowed if at least one *Allow* statement  
 181 matches the request, and none of the *Deny*  
 182 statements match the request.

183 Specifically, *Effect* specifies whether  
 184 the request is allowed or denied. *Action*  
 185 and *Resource* define the permitted or pro-  
 186 hibited actions and cloud resources (e.g.,  
 187 S3 buckets, IAM roles). *Principal* indi-  
 188 cates the entities performing the actions.

189 *Condition* imposes additional constraints (e.g., IP address or time). A *Condition* is defined as a list  
 190 of expressions. Each expression applies an operator (*OpName*) to one or more values (*Value*) for a  
 191 condition key (*KeyName*), where the types of values include strings, bit-vectors, ranges, enumera-  
 192 tions, etc. The operator determines both the type of the values and the semantics of the comparison.  
 193 For example, *StringEquals* checks whether two strings are identical, whereas *StringLike* checks

```

Policy → [Statement]
Statement → (Effect, Principal, Action, Resource,
Condition?)
Effect → Allow | Deny
Principal → Principal : [string]
Action → Action : [string]
Resource → Resource : [string]
Condition → Condition : [Operator]
Operator → (OpName, KeyName, [Value])
OpName → StringEquals | StringLike | IPAddress | ...
KeyName → sourceVpc | sourceIP | prefix
Value → string | bitvector | range | enum | ...

```

Fig. 1. Simplified abstract syntax for the AWS policy language, where "?" denotes optional elements and "[]" denotes a list of valued elements.

```

Effect      : "Allow"    // Statement1
Principal   : "*"
Resource    : "dept*/user1.txt",
              "dept1/user*.txt"
IpAddress  : "112.0.0.0/24", "113.0.0.0/24"
-----
Effect      : "Deny"     // Statement2
Principal   : "*"
NotResource : "dept*/user1.txt"
IpAddress  : "112.0.0.0/24"
-----
Effect      : "Deny"     // Statement3
Principal   : "*"
NotResource : "dept1/user*.txt"
IpAddress  : "113.0.0.0/24"

```

Fig. 2. An example of an *IAM policy* with a simplified for-  
 mat. For clarity, we omit some necessary keys, such as *Action*, and omit the condition key for *IpAddress*.

197 whether a string matches a pattern including wildcards. Moreover, the `IpAddress` checks whether  
 198 an IP address (i.e., a bit vector) is within an IP prefix.

199 To achieve flexible and complex access control, *IAM policies* rely on two mechanisms. We take  
 200 Fig. 2 as an example.

201 (i) *Wildcard semantics*. String operators provide pattern matching, where `*` matches any number  
 202 of characters, and `?` matches a single character (equivalent to `"*"` and `"?"` in regular expressions, re-  
 203 spectively). For example, `dept1/user*.txt` matches any file in the `dept1` directory whose name  
 204 starts with `user` and ends with `.txt`, such as `user123.txt`. `IpAddress` also supports pattern  
 205 matching. It represents a range of IP addresses using CIDR notation, where a prefix in the for-  
 206 mat `"A/B"` specifies all IP addresses that share the first B bits with IP address A. For example,  
 207 `112.0.0.0/24` restricts the range from `112.0.0.0` to `112.0.0.255`.

208 (ii) *Boolean semantics*. With one domain (key), different values together represent an OR rela-  
 209 tionship. Between different keys, the relationship is AND. NOT can be indicated by adding `Not`  
 210 before the domain, such as `NotResource`.

211 In the example of Fig. 2, the semantics of Statement2 are as follows: any request whose Re-  
 212 source does **not** match `"dept*/user1.txt"` and whose `IpAddress` is within `"112.0.0.0/24"`  
 213 will be denied. For example, the following request is denied by Statement2: the file path `"dept1/user2.txt"`  
 214 does not match pattern `"dept*/user1.txt"`, and IP address `"112.0.0.32"` falls within IP prefix  
 215 `"112.0.0.0/24"`.

216 `Principal : "user1" Resource : "dept1/user2.txt" IPAddress : "112.0.0.32"`

217 The semantics of the policy are as follows: a request is allowed if it matches Statement1 but  
 218 does not match Statement2 or Statement3; otherwise, it is denied. Thus, the following request is  
 219 allowed by the policy in Fig. 2 because it satisfies this condition.

220 `Principal : "user1" Resource : "dept1/user1.txt" IPAddress : "112.0.0.32"`

221 This example illustrates that determining which requests are allowed and which are denied by a  
 222 policy can be challenging. In practice, real-world policies are even more complex, making it crucial  
 223 to identify the requests that a policy permits.

## 226 2.2 Access Control Model

227 Our access control model is identical to the model of AWS Access Analyzer[4]. The difference is  
 228 that we define stricter conciseness of the intents, as defined in §4.2.

229 **Requests.** Let  $K = \{k_1, \dots, k_n\}$  be a set of keys. Let  $V_k = \{v_1, \dots\}$  be a (possibly infinite) set of values  
 230 for key  $k$ . A *request*  $r$  is a mapping from keys  $k$  to values in  $V_k$ . For example,  $r = \{\text{Principal} \mapsto$   
 231 `"user1"`, `Resource`  $\mapsto$  `"dept1/user2.txt"`, `IpAddress`  $\mapsto$  `"112.0.0.32"` $\}$ .

232 **Policies.** A *policy* is a predicate on requests  $\mathbb{P} : r \rightarrow \text{Bool}$ . The allowed requests of a policy  $\mathbb{P}$  is  
 233 defined as:

$$\sigma(\mathbb{P}) = \{r \mid \mathbb{P}(r) = \text{true}\} \quad (1)$$

234 **Labels.** A *label*  $l_k$  is a set of values for a key  $k$ , expressed in a human-readable format. All labels  $l_k$   
 235 for a key  $k$  form the universal set  $L_k$ . For example, for the key `IpAddress`, the universal set of labels  
 236 is  $L_{\text{IpAddress}} = \{0.0.0.0/24, 112.0.0.0/24, 113.0.0.0/24\}$ . Similarly, for the key `Resource`, the  
 237 universal set of labels is  $L_{\text{Resource}} = \{"*", "dept1/user*.txt", "dept*/user1.txt"\}$ .

238 **Intents.** A intent  $\mathbb{I}$  is a map from keys  $K$  to labels  $L$ . The allowed request of an intent  $\mathbb{I}$  is the set  
 239 of requests where each key  $k$  is mapped to a value in the label  $\mathbb{I}(k)$ :

$$\sigma(\mathbb{I}) = \{r \mid \forall k. r(k) \in \mathbb{I}(k)\} \quad (2)$$




Fig. 3. The processes of *Access Analyzer* for Fig. 2. For simplicity, we only consider two keys, i.e., Resource and IpAddress.

Let  $\mathbb{S}$  be a set of intents, i.e.,  $\mathbb{S} = \{I_1, I_2, \dots, I_n\}$ . Then, the allowed requests of  $\mathbb{S}$  is defined as:

$$\sigma(\mathbb{S}) = \bigcup_{I \in \mathbb{S}} \sigma(I) \quad (3)$$

For example, the policy in Fig. 2 has the following two intents:

Principal : "\*" Resource : "dept\*/user1.txt" IpAddress : "112.0.0.0/24" //Intent1  
Principal : "\*" Resource : "dept1/user\*.txt" IpAddress : "113.0.0.0/24" //Intent2

As we can see, the two intents only have positive (allow) declarations, without any negative (deny) declarations. This is to ensure the intents are easier to understand. It is easy to see this set of intents is *sound*, as it includes all requests permitted by the policy. It is also *precise*, containing minimal requests outside the policy (in this example, none). Furthermore, this set of intents is *concise*, since removing any intent results in incomplete coverage of the policy. To formalize these goals, we introduce a set of formal properties (see §4.2).

### 2.3 Limitations of Existing Approaches

Existing methods for analyzing *IAM policies*, including intent mining, are based on SMT (Satisfiability Modulo Theories). These methods encode both policies and intents [4, 5, 8, 16, 17, 31, 41] as SMT constraints. Typically, a policy is modeled as a constraint  $\mathbb{P} = \bigvee_{S \in \text{Allow}} S \wedge \neg \bigvee_{S \in \text{Deny}} S$ , where each statement  $S$  is encoded as an SMT constraint.

To the best of our knowledge, the *Access Analyzer* deployed by AWS is the only state-of-the-art intent mining tool documented in the literature[4]. Some tools, such as AWS Zelkova[5] and Microsoft Ambit[31], only determine the implication relationships between two policies. However, even though *Access Analyzer* can ensure the soundness and precision requirements for intent mining, it still cannot meet the other two requirements, i.e., speed and conciseness.

**Limitation in speed due to the redundancy among multi-round SMT solving.** Mining intents [4] invokes the SMT solver independently for each intent, resulting in a large number of solver invocations. For example, *Access Analyzer* adopts a stratified refinement process based on predefined *labels*, as illustrated in Fig. 3. Starting from an over-permissive intent  $I_1(*,*)$  (ensuring *soundness*), it iteratively checks the following constraint<sup>1</sup>. Here,  $\text{Child}(I)$  (defined in Formula 8) denotes the set of intents that are maximally covered by the current intent (e.g.,  $I_2, I_3, I_4, I_5$ ).

$$\mathbb{P} \wedge \left( I_i \wedge \neg \left( \bigvee_{I_j \in \text{Child}(I_i)} I_j \right) \right) \quad (4)$$

<sup>1</sup>To enable efficient solving, [4] uses a simplified but equivalent version of this constraint. For clarity, we present the original form.

If the constraint is satisfiable, the intent is included in the final set of intents, indicating that it captures a unique part of the policy (ensuring *precision*); otherwise, its child intents are added to the worklist, and the above process repeats until the worklist is empty. As shown in Fig. 3, 9 intents are checked. Moreover, even a policy with only 15 statements can trigger  $O(10^3)$  SMT solver (§6.6) invocations.

Meanwhile, SMT solving per query is expensive due to the iterative process between theory solvers and the SAT solver. As shown in Fig. 4, after the SMT constraint is converted into conjunctive normal form (CNF) [45] and theory-specific parts are replaced with Boolean variables [10], a new constraint  $\Phi$  is produced. An SAT solver searches for a satisfying assignment  $M$  for  $\Phi$ , which is then iteratively checked by theory solvers. If a theory solver check fails (T-UNSAT), the solver backtracks and tries a new assignment. This loop continues until  $M$  is T-SAT or  $\Phi$  is UNSAT[23]. Thus, each round of SMT solving is expensive, ranging from 28ms to 4545ms (see § 6.6), causing the mining process to take over an hour (§6.3).

However, checking whether a policy and an intent overlap requires full SMT solving, involving multiple theories such as string theory, bit-vector theory, and so on. SMT constraints cannot be simplified through syntax-level transformations.

**Limitation in conciseness due to not considering cross-intent relations.** Fig. 3 shows that four intents,  $I_6, I_7, I_8, I_9$ , are generated by *Access Analyzer*, as their corresponding constraints (Formula 4) are satisfiable. However, §2.1 demonstrates that only two intents,  $I_6$  and  $I_9$ , are sufficient. This redundancy arises because the four intents partially overlap with one another but are not fully contained within each other. However, these intents still contain redundancy, i.e.,  $(I_6 \cup I_7 \cup I_8 \cup I_9) \cap P \subseteq (I_6 \cup I_9) \cap P$ , where  $P$  (or  $I$ ) denotes the set of permitted requests. Meanwhile, *Access Analyzer* does not detect such relationships. As a result, a policy with only 10 statements can yield results containing more than 100 intents, which is  $10\times$  larger than the original policy (§6.2).

However, selecting a minimal subset of intents from the candidate set requires multi-round SMT solving again. Given a candidate set of intents  $\mathbb{S}$ , the search space grows exponentially as  $O(2^{|\mathbb{S}|})$ . Each enumeration requires the SMT solver to check whether the selected intents together cover the policy.

**Existing works hard to address the limitations.** Quacky [17] employs the automata-based solver ABC [3] for model counting, i.e., enumerating the number of policy-allowed requests, but it is less efficient than SMT solvers for single-round satisfiability checks. Similarly, [13] applies the equivalence class approach to count the number of IPs permitted by a policy, but it only supports only IP and lacks regular expression semantics. In addition, Margrave [21] leverages BDDs [2] to check implication between two policies under XACML (rather than IAM) semantics, supporting only exact string matching. Other tools [5, 8, 31, 41] similarly focus on analyzing a single policy or checking implication between two policies. However, all of these approaches do not support accelerating multi-round solving, and thus are difficult to apply to intent mining in IAM policies.




Fig. 4. The basic idea of DPLL(T).




Fig. 5. The example of Fig. 2. For clarity, we only consider two domains Resource, IpAddress.

### 3 Overview

In this section, we present two key insights: (1) accelerating intent mining by pre-extracting theory redundancy across multi-round SMT solving, and (2) selecting minimal intents by reducing the problem to a *min-set-cover* over finite sets. We then describe the overall workflow.

#### 3.1 Key Insights


**Insight 1: The redundancy for theory solvers can be pre-extracted, thereby avoiding theory solving in multi-round SMT solving.** For intent mining, we observe that even though it involves multi-round SMT solving, the possible values of variables are known in advance. Taking Fig. 2 as an example, the value of Resource is selected from the set  $L_{\text{Resource}} = \{\ast, \text{dept}1/\text{user}.txt, \text{dept}^*/\text{user}1.txt\}$ . In another case, the value of IpAddress comes from the set  $L_{\text{IpAddress}} = \{0.0.0.0/24, 112.0.0.0/24, 113.0.0.0/24\}$ . The difference in multi-round SMT solving lies in how these possible values are combined.

Since all possible values of all constraints are known in advance, this enables us to partition these possible values (i.e., *labels*) into disjoint spaces, i.e., *equivalence classes (ECs)* [49], where two values in the same EC yield the same satisfiability result for the theory solver. Returning to the example in Fig. 2, the *label* set  $L_{\text{Resource}}$  can be partitioned into four ECs  $\boxed{0}, \boxed{1}, \boxed{2}, \boxed{3}$ , as shown in Fig. 5 (a). The *label* set  $L_{\text{IpAddress}}$  can be partitioned into three ECs  $\boxed{0}, \boxed{1}, \boxed{2}$  as shown in Fig. 5 (b).

Then, we can maintain a mapping  $\mathcal{M}$  from each *label* to its corresponding ECs, e.g.,

- $\text{dept}^*/\text{user}1 \mapsto \{\boxed{0}, \boxed{1}\}$
- $\text{dept}1/\text{user}^* \mapsto \{\boxed{0}, \boxed{2}\}$
- $\ast \mapsto \{\boxed{0}, \boxed{1}, \boxed{2}, \boxed{3}\}$

Based on the mapping, we can quickly obtain single-theory constraints over bit-vectors that are equivalent to multi-theory SMT constraints. We use Statement1 and Statement2 in Fig. 2 as illustrative examples. We first obtain their ECs from the mapping. As shown in Fig. 5(c), the ECs of Statement1 in the Resource domain are  $\mathcal{M}(\text{dept}^*/\text{user}1) \cup \mathcal{M}(\text{dept}1/\text{user}^*) = \{\boxed{0}, \boxed{1}, \boxed{2}\}$  (recall that, for one key, different values together represent an OR relationship). Similarly, the ECs in the IpAddress domain are  $\mathcal{M}(112.0.0.0/24) \cup \mathcal{M}(113.0.0.0/24) = \{\boxed{0}, \boxed{1}\}$ . Thus, Statement1 can be represented as the Cartesian product  $\{\boxed{0}, \boxed{1}, \boxed{2}\} \times \{\boxed{0}, \boxed{1}\}$  (recall that, for different keys, the relationship is AND). We then encode it as a bit-vector, as shown on the left of Fig. 5(c). There is a special case in the figure where  $\mathcal{M}(\neg \text{dept}^*/\text{user}1) = \{\boxed{0}, \boxed{1}, \boxed{2}, \boxed{3}\} - \{\boxed{0}, \boxed{1}\} = \{\boxed{2}, \boxed{3}\}$  (recall that, NOT can be indicated by adding Not before the domain). After encoding each statement, the policy  $\mathbb{P}$  can be expressed as  $\mathbb{P} = \bigvee_{S \in \text{Allow}} S \wedge \neg \bigvee_{S \in \text{Deny}} S$ . In this way, both the policy and the intents are encoded using only bit-vector constraints.



400 Fig. 6. An example of encoding the selection of minimal intents as *min-set-cover* problem.

402 Note the pre-computing of equivalence classes is performed only once, and the results are stored  
403 in the mapping. In this way, we extract the redundancy among multiple rounds of invoking theory  
404 solvers, as shown in Fig. 4. We will show our approach speeds up intent mining by one to two  
405 orders of magnitude compared with SMT solvers (§6.3).

406 **Insight 2: Reducing intents requires covering all allowed requests of the policy, and can  
407 then be formulated as a min-set-cover problem.** We observe that to guarantee *soundness*, the  
408 space of requests specified by the intents must fully cover the space of requests allowed by the  
409 policy. Therefore, the problem of minimizing the number of intents can be naturally formulated  
410 as a *min-set-cover* problem: *selecting the minimal intents from candidate intents that still cover all  
411 requests of the policy*, which ensures no valid requests are omitted.

412 Some classic works [9, 26, 48] address the *min-set-cover* problem over finite sets. However, both  
413 policy and intents can potentially span infinite spaces (e.g., dept\*/user1.txt matches infinitely  
414 requests). Therefore, we employ equivalence class (EC) techniques to transform the *min-set-cover*  
415 problem over infinite sets into an equivalent problem over finite sets, thereby enabling the use of  
416 existing algorithms for the *min-set-cover* problem. We take the policy in Fig. 2 and the candidate  
417 intents  $\mathbb{I}_6$ ,  $\mathbb{I}_7$ ,  $\mathbb{I}_8$ , and  $\mathbb{I}_9$  in Fig. 3 as the example. The policy space and intents space (both infinite)  
418 are shown in Fig. 6(a) and (b). The request space, over both the policy space and the intents space,  
419 is partitioned into six disjoint spaces (ECs), i.e.,  $\{a, b, c, d, e, f\}$ , as shown in Fig. 6(c). Under this  
420 partitioning,  $\mathbb{I}_6 = \{d, e\}$ ,  $\mathbb{I}_7 = \{a, b\}$ ,  $\mathbb{I}_8 = \{d, f\}$ , and  $\mathbb{I}_9 = \{a, c\}$ , while  $\mathbb{P}$  corresponds to  $\{a, c, d, e\}$ .  
421 Therefore, the problem reduces to selecting intents that cover  $a, c, d, e$  in  $\mathbb{P}$ . As a result,  $\mathbb{I}_6$  and  $\mathbb{I}_9$   
422 are selected.

423 However, the existing tools compute *equivalence classes* (ECs) primarily using an iterative method  
424 [33, 34, 49, 51]. For example, we consider  $\mathbb{P}$  and  $\mathbb{I}_6$ , resulting in three ECs:  $\neg\mathbb{I}_6 \wedge \mathbb{P}$ ,  $\mathbb{I}_6 \wedge \neg\mathbb{P}$ ,  $\mathbb{I}_6 \wedge \mathbb{P}$ .  
425 Among these,  $\mathbb{I}_6 \wedge \neg\mathbb{P}$  is empty. Next, we consider  $\mathbb{I}_7$  and compute intersections with the remaining  
426 ECs:  $\neg\mathbb{I}_7 \wedge (\neg\mathbb{I}_6 \wedge \mathbb{P})$ ,  $\mathbb{I}_7 \wedge (\neg\mathbb{I}_6 \wedge \mathbb{P})$ , ... . We find that even after converting the SMT constraints  
427 of policy and intent into bit-vector constraints and employing an SAT solver to check the satis-  
428 fiability of each constraint, the search space remains large due to the involvement of up to 100  
429 intents.

430 Therefore, we adopt *Binary Decision Diagrams* (BDDs) to compactly represent bit-vectors instead  
431 of using the SAT solver. **The insight is that BDDs allow more efficient re-use of intermediate results  
432 for incremental computation.** BDD is a directed acyclic graph (DAG) with terminal nodes (true  
433 and false) and variable nodes. Each variable node assigns a Boolean variable  $x$ , with two outgoing  
434 edges: a solid line and a dashed line, representing true and false assignments, respectively. Each  
435 path from the root to the “true” terminal node represents a truth assignment of the constraint.  
436 For clarity, we illustrate the benefits of BDDs with a simple example. Consider two formulas: (1)  
437  $11** \wedge 1*0*$  and (2)  $11** \wedge 1*0* \wedge *100$ . After computing formula (1) with BDDs, we obtain a BDD  
438 representing  $110*$ . When computing formula (2), we can reuse the existing BDD and incrementally  
439 compute  $110* \wedge *100$ . In contrast, although SAT solvers support push/pop operations, they still  
440

need to consider the full formula (2) and only leverage learned clauses to accelerate the search, which is generally less efficient than BDDs. We now consider the scenario of computing ECs.

Fig. 7 shows the BDDs of  $\mathbb{P}$ ,  $\mathbb{I}_6$  and  $\neg\mathbb{I}_6 \wedge \mathbb{P}$  respectively. When computing  $\mathbb{I}_7 \wedge (\neg\mathbb{I}_6 \wedge \mathbb{P})$ , BDD avoids recomputing  $\neg\mathbb{I}_6 \wedge \mathbb{P}$ . Instead, BDD compactly represents the solution space of  $\neg\mathbb{I}_6 \wedge \mathbb{P}$  so that  $\mathbb{I}_7$  operates directly on it.

We will show that our approach reduces the number of intents by one order of magnitude (§6.2) and speeds up intent reduction by one to four orders of magnitude compared with SMT solvers (§6.3). Moreover, in our approach, using BDD to represent bit-vectors is one to two orders of magnitude faster than using SAT solver (§6.5).




Fig. 7. The *binary decision diagram (BDD)* for  $\mathbb{P}$ ,  $\mathbb{I}_6$  and  $\neg\mathbb{I}_6 \wedge \mathbb{P}$ , respectively.

### 3.2 Workflow of AccessRefinery

*AccessRefinery* mines the intents of an *IAM policy*, as shown in Fig. 8. It takes the *IAM policy* as input and produces the corresponding set of intents. To achieve this, *AccessRefinery* consists of three major modules: the *Intent Miner*, the *Intent Reducer*, and the *Multi-theory Constraint Preprocessor* (MCP). Among them, MCP decouples the underlying Boolean operations from the higher-level application logic, which abstracts the low-level operations and allows the upper-layer algorithms to focus solely on their logic. *AccessRefinery* is fully automatic, requiring no operator intervention, and its workflow is as follows:


*AccessRefinery* first initializes MCP, which (1) automatically collects all variables according to their type, e.g., regular expressions and IP prefixes; (2) maps them to the corresponding operable data structures, e.g., automata and BDDs; (3) maps each operable data structure (e.g., automata and BDDs) to a unified superclass based on language features, e.g., polymorphism in Java; and (4) consistently applies the EC partitioning algorithm.

Then, *Intent Miner* generates a set of raw (non-reduced) intents based on the input policy. It adopts the existing stratified framework [4], but instead of directly using SMT solvers, it relies on MCP for faster multi-round SMT solving.

Finally, *Intent Reducer* generates a minimum set of intents from the set of raw intents, such that they can still cover all raw intents. *Intent Reducer* achieves this by modeling a min-set-cover problem, which MCP transforms into a min-set-cover problem over a finite set, and solving it with an Integer Linear Programming (ILP) solver [24].

We highlight that MCP differs from existing approaches by employing a two-level encoding scheme. First, MCP encodes each domain label using a data structure appropriate for its type: strings are represented using regular expressions, prefixes using BDDs, ranges using red-black trees, and enumerated variables using sets (first level). MCP then partitions equivalence classes and encodes them into BitVectors. These BitVectors are subsequently represented as BDDs (second level), and all SMT solving is performed as Boolean operations on this second-level BDD.

In contrast, prior work [13] directly partitions equivalence classes for IP addresses and then performs single-domain operations on the set of equivalence classes for IP prefixes. Margrave [21] maps strings directly to BDDs by encoding each character as a separate BDD variable; the entire string is represented as the conjunction of these variables. All subsequent operations are performed directly on this BDD, which supports only exact string matching and does not support regular expressions or multi-round SMT solving for acceleration.

Fig. 8. The workflow of *AccessRefinery*.

## 4 Design Details

### 4.1 Multi-Theory Constraint Preprocessor (MCP)

Given the fundamental role of Multi-Theory Constraint Preprocessor (MCP) in both intent mining and reduction, we first introduce the design of MCP.

**The APIs.** MCP hides the details of calculating ECs for each domain and encoding them as BDDs, allowing upper-layer applications to focus on their own logic. For example, Fig. 9 illustrates the use of MCP to solve  $\neg \mathbb{I}_6 \wedge \mathbb{I}_7$  without exposing the underlying details. The APIs of MCP are explained as follows:

- `AddValue(Domain, Type, Value)` defines a variable in Domain with the specified Type and Value. For IAM policies, supported types include REGEXP (e.g., the Resource key), PREFIX (e.g., the IpAddress key), RANGE (e.g., the DateGreaterThan key), and ENUM (e.g., the Action key).
- `PartitionECs()` partitions all values into ECs for each domain and maps each *label* to its corresponding ECs. This function is called after all domain values have been declared.
- `GetValue(Domain, Value)` returns the BDD representation of a Value in a Domain. SMT solving over *labels* is equivalently converted into BDD operations.

**Definition of Equivalence Classes.** A *label* is denoted as  $l = \text{false}$  if it is a contradiction, e.g.,  $l = 112.0.0.0/24 \wedge 113.0.0.0/24$ , and as  $l = \text{true}$  if it is a tautology, e.g.,  $l = 0.0.0.0/0$ .

Following [49], we define ECs as follows:

DEFINITION 1. Given a set of labels  $\mathcal{L}$ , a set of equivalence classes  $\mathcal{C} = \{c_1, \dots, c_N\}$  satisfies:

- $c_i \neq \text{false}, \forall i \in \{1, \dots, N\}$ ,
- $\bigvee_{i=1}^N c_i = \text{true}$ ,
- $c_i \wedge c_j = \text{false}$ , if  $i \neq j$ ,
- Each  $l \in \mathcal{L}, l \neq \text{false}$ , can be expressed as the disjunction of a subset of equivalence classes:

$$l = \bigvee_{i \in \mathcal{M}(l)} c_i, \quad \mathcal{M}(l) \subseteq \{1, \dots, N\},$$

- $N$  is the minimal number that the above conditions hold.

The above definition induces a mapping  $\mathcal{M} : \mathcal{L} \rightarrow 2^{\{1, \dots, N\}}$  from *label* to its corresponding ECs. For example,  $\mathcal{M}(\text{dept*/user1}) = \{\text{①}, \text{②}\}$ . This decomposition is *unique* for any  $l \in \mathcal{L}$  [49]. Crucially, logical operations over *labels* are reduced to efficient set operations:  $\mathcal{M}(l_1 \wedge l_2) = \mathcal{M}(l_1) \cap \mathcal{M}(l_2)$ ,  $\mathcal{M}(l_1 \vee l_2) = \mathcal{M}(l_1) \cup \mathcal{M}(l_2)$  [49].

**Computation of Equivalence Classes.** After collecting all values for each domain, we calculate the EC and mapping  $\mathcal{M}$  as follows:

*Step (1) Domain-Specific Label Representation.* Partitioning ECs requires Boolean operations, e.g., conjunction ( $\wedge$ ) and negation ( $\neg$ ). However, *labels* cannot directly perform this operation. For example,  $113.0.0.0/24 \wedge \neg 112.0.0.0/24$  cannot be represented as a single CIDR block. Therefore, we map each *label* to a corresponding data structure that supports equivalent Boolean operations.

```

540  /* Declare the domain values. */
541  MCP mcp = new MCPFactory();
542  mcp.addValue("Res", REGEXP, "dept*/user1.txt");
543  mcp.addValue("Res", REGEXP, "dept1/user*.txt");
544  mcp.addValue("IP", PREFIX, "112.0.0.0/24");
545  mcp.addValue("IP", PREFIX, "113.0.0.0/24");
546  /* Extract theory redundancy.*/
547  mcp.partitionECs();
548  /* Multi-round SMT solving.*/
549  BDD res1 = mcp.getValue("Res", "dept*/user1.txt");
550  BDD res2 = mcp.getValue("Res", "dept1/user*.txt");
551  BDD ip1 = mcp.getValue("IP", "112.0.0.0/24");
552  BDD ip2 = mcp.getValue("IP", "113.0.0.0/24");
553  BDD s1 = (res1.or(res2)).and(ip1.or(ip2));
554  BDD s2 = res1.not().and(ip1);
555  BDD s3 = res2.not().and(ip2);
556  BDD policy = s1.diff(s2).diff(s3);
557  BDD intent6 = res1.and(ip1);
558  if (policy.and(intent6.not()).sat() == true) {...}
559

```

Fig. 9. An example usage of MCP.

- Regexp *labels* are mapped to automatons [27]. An automaton is a finite-state machine capable of recognizing patterns described by regular expressions. For example, `dept*/user1.txt`  $\vee$  `dept*/user2.txt` is represented as an automaton that accepts both patterns.
- BitVec *labels* are mapped to nodes in a Binary Decision Diagram (BDD) [2], a data structure used to represent Boolean functions. For example, `112.0.0.0/24`  $\wedge \neg$  `112.0.0.128/25` results in a BDD representing the IP range `112.0.0.0/25`.
- Range *labels* are mapped to nodes in a balanced binary search tree (BST) [28]. For example, `[9, 18]  $\vee$  [19, 24]` results in two disjoint intervals, represented as two nodes in the tree.
- Enum *labels* are mapped to sets, i.e., collections of unique elements. For example, `List-Bucket`  $\vee$  `CreateBucket` results in a set with two elements.

*Step (2) Domain-Independent ECs Partitioning.* For a given set of *labels*  $\mathcal{L}$ , Algorithm 1 generates the set of ECs  $\mathcal{C}$  and the mapping  $\mathcal{M}$ . Specifically, Line 2 initializes the current set of ECs with the *true label*. In Line 3, it iterates over each *label*  $l \in \mathcal{L}$ . For each *label*  $l$  and each current EC  $c \in \mathcal{C}$ , it computes the conjunctions  $l \wedge c$  and  $l \wedge \neg c$  (Line 6). Only non-empty results are included in the current EC set. Line 7 updates  $\mathcal{C}$  with the new ECs for the next iteration. Lines 8–10 construct the mapping  $\mathcal{M}$  by associating each *label*  $l$  with the set of ECs. Finally,  $\mathcal{C}$  and  $\mathcal{M}$  are returned.

Many prior approaches [13, 49, 51] employ a similar approach to partition equivalence classes (ECs). However, they are limited to EC partitioning over IP prefixes. Our approach differs from previous work in that we use labels representing abstract semantics. By mapping variables of different types to corresponding operable data structures, we then perform EC partitioning on a unified abstract representation.

**Label Encoding.** After calculating the mappings from *label* to ECs, we encode SMT constraints over *labels* as follows:

*Step (1) Encoding EC as Bit-Vectors.* Suppose there are  $m$  domains, and the  $k$ -th domain contains  $N_k$  ECs. We adopt binary encoding, where the total number of bits required is  $\sum_{k=1}^m \lceil \log_2(N_k) \rceil$ . Each domain is assigned a unique range  $\text{Position}(k) = \left[ \sum_{j=1}^{k-1} \lceil \log_2(N_j) \rceil + 1, \sum_{j=1}^k \lceil \log_2(N_j) \rceil \right]$ . We next define  $\mathcal{V}(i, k)$  as a bit-vector encoding the  $i$ -th EC in the  $k$ -th domain. Let  $\text{binary}(a, b)$

---

**Algorithm 1:** Compute ECs for a single domain
 

---

```

589 Input :  $\mathcal{L}$ : the set of labels.
590 Output :  $\mathcal{C}$ : the set of equivalence classes for  $\mathcal{L}$ .  $\mathcal{M}$ : the mapping from label to its ECs.
591
592 1 Function ComputeSECs( $\mathcal{L}$ ):
593   2    $\mathcal{C} \leftarrow \{\text{true}\}$ 
594   3   foreach  $l \in \mathcal{L}$  do
595     4      $\mathcal{C}' \leftarrow \{\}$ 
596     5     foreach  $c \in \mathcal{C}$  do
597       6        $\mathcal{C}' \leftarrow \mathcal{C}' \cup \{l \wedge c, l \wedge \neg c\}$             $\triangleright$  calculating equivalence classes
598     7      $\mathcal{C} \leftarrow \mathcal{C}'$ 
599
600   8   foreach  $l \in \mathcal{L}, c \in \mathcal{C}$  do
601     9     if  $l \wedge c \neq \text{false}$  then
602       10        $\mathcal{M}(l) \leftarrow \mathcal{M}(l) \cup \{id(c)\}$             $\triangleright$  calculating the mapping from label to its ECs
603
604 11 return  $\mathcal{C}, \mathcal{M}$ 
605
606

```

---

denote the bit-vector representing the range from integer  $a$  to  $b$ . For example,  $\text{binary}(0, 2) = 0* \vee 10$ . Then,  $\mathcal{V}(i, k)$  assigns  $\text{binary}(i, i)$  to the bit range specified by  $\text{Position}(k)$ . For all other domains  $j \neq k$ , the bits in their respective ranges  $\text{Position}(j)$  are set to  $\text{binary}(0, N_j - 1)$ , where  $N_j$  is the number of ECs in domain  $j$ .

*Step (2) Encoding labels as BDDs.* Each *label*  $l$  belongs to a specific domain  $k$  and has corresponding ECs  $\mathcal{M}_k(l)$ . Let  $\mathcal{B}(v)$  return the BDD for the bit-vector  $v$ . Then, the BDD encoding of *label*  $l$  is  $\mathcal{F}(l) = \bigvee_{i \in \mathcal{M}_k(l)} \mathcal{B}(\mathcal{V}(i, k))$ .

**THEOREM 1.** Let  $l_1$  and  $l_2$  be labels, and let  $\circ \in \{\wedge, \vee, \setminus, \Rightarrow, \Leftrightarrow\}$  be a Boolean operation. Then,  $\mathcal{F}(l_1 \circ l_2) = \mathcal{F}(l_1) \circ \mathcal{F}(l_2)$  and  $\mathcal{F}(\neg l_1) = \neg \mathcal{F}(l_1)$  both hold.

**PROOF SKETCH.** Mapping each domain label to its corresponding operable data structure preserves the equivalence of Boolean operations, as established in [2, 32, 38, 43] for Regexp, BitVec, Range, and Enum, respectively. Furthermore, encoding equivalence classes (ECs) as bitvectors also preserves the equivalence, since each domain is assigned a non-overlapping segment of the bitvector, ensuring a one-to-one correspondence between the ECs and their bitvector representations.  $\square$

## 4.2 Intent Miner

In this section, we show the goals that high-quality intents should satisfy, and then describe how to generate such intents.

**Goals.** To satisfy soundness, precision, and conciseness, we define the following properties, inspired by [4]. Recall that  $\sigma(\mathbb{P}), \sigma(\mathbb{I}), \sigma(\mathbb{S})$  denote the set of allowed requests of the policy, intent, and set of intents, respectively.

- **Soundness (Coverage):** The set of intents must cover all allowed requests:

$$\sigma(\mathbb{P}) \subseteq \sigma(\mathbb{S}) \tag{5}$$

- **Precision (Irreducibility):** Each intent captures a unique part of the policy and cannot be further refined:

$$\forall \mathbb{I} \in \mathbb{S}, \exists r \in \sigma(\mathbb{I}) \cap \sigma(\mathbb{P}), \forall \sigma(\mathbb{I}') \subseteq \sigma(\mathbb{I}), r \notin \sigma(\mathbb{I}') \tag{6}$$

- **Conciseness (Minimality):** No intent is redundant with respect to the others:

$$\forall \mathbb{I} \in \mathbb{S}, \sigma(\mathbb{I}) \cap \sigma(\mathbb{P}) \not\subseteq \bigcup_{\mathbb{I}' \in \mathbb{S}, \mathbb{I} \neq \mathbb{I}'} \sigma(\mathbb{I}') \cap \sigma(\mathbb{P}) \quad (7)$$

We highlight that our definition of *minimality* is stricter than in prior work [4], which only analyzes the mutual inclusion relationship between two intents.

**Intent Mining.** *Intent Miner* adopts a stratified analysis approach to ensure both *coverage* and *irreducibility*. While this approach is inspired by *Access Analyzer* [4], we replace its underlying SMT solver with our customized MCP engine. Specifically, *Intent Miner* initializes a *worklist* set containing a single element, i.e., the universe intent, and an empty *result* set. We then iteratively process the *worklist*: for each current intent, we remove it from the *worklist* and enumerate all its *maximal subsets* via  $\text{Child}(\mathbb{I})$ . These maximal subsets are defined as intents included by the current intent and not implied by any other intent, where  $\mathbb{I}' \prec \mathbb{I}$  means  $\sigma(\mathbb{I}') \subseteq \sigma(\mathbb{I})$ .

$$\text{Child}(\mathbb{I}) \triangleq \{\mathbb{I}' \mid \mathbb{I}' \prec \mathbb{I} \wedge \forall \mathbb{I}''. \neg(\mathbb{I}' \prec \mathbb{I}'' \prec \mathbb{I})\} \quad (8)$$

Next, *Intent Miner* checks Formula 4. If Formula 4 is satisfiable, we add it to the *result*. If the formula is unsatisfiable, we add all proper subsets of the current intent to the *worklist*. This process repeats until the *worklist* becomes empty. According to [4], the result satisfies both *coverage* and *irreducibility*.

### 4.3 Intent Reducer

In this section, we first formulate the task of selecting a minimal subset of intents as a *min-set-cover* problem. We then describe how this problem is reduced to a finite instance using ECs, and finally solved using an Integer Linear Programming (ILP) solver[24].

**Problem Formulation.** Let the set of candidate intents be  $\mathbb{S} = \{\mathbb{I}_1, \mathbb{I}_2, \dots, \mathbb{I}_n\}$ , generated by *Intent Miner*. These candidate intents satisfy both *coverage* and *irreducibility* [4].

However, the number of intents is often large, motivating the need to select a smaller subset. Removing intents may lead to a loss of *coverage*, as some requests allowed by the policy might no longer be covered. Fortunately, *irreducibility* is preserved under subset selection:

LEMMA 1. *If  $\mathbb{S}$  is irreducible, then any non-empty subset  $\mathbb{S}' \subseteq \mathbb{S}$  is also irreducible.*

PROOF SKETCH. We prove it by contradiction. Assume that a non-empty subset  $\mathbb{S}' \subseteq \mathbb{S}$  is *reducible*. Then, (1)  $\exists \mathbb{I}' \in \mathbb{S}'$  such that  $\forall r \in \sigma(\mathbb{I}') \cap \sigma(\mathbb{P}), \exists \sigma(\mathbb{I}'') \subseteq \sigma(\mathbb{I}'), r \in \sigma(\mathbb{I}'')$ . However, since  $\mathbb{S}$  is *irreducible* and  $\mathbb{I}' \in \mathbb{S}$  (because  $\mathbb{S}' \subseteq \mathbb{S}$ ), the *irreducibility* condition demands that (2)  $\exists r \in \sigma(\mathbb{I}') \cap \sigma(\mathbb{P})$  such that  $\forall \sigma(\mathbb{I}'') \subseteq \sigma(\mathbb{I}'), r \notin \sigma(\mathbb{I}'')$ . This leads to a contradiction between (1) and (2). Hence, all non-empty subsets of  $\mathbb{S}$  must be *irreducible*.  $\square$

We now formally state our goal:

PROBLEM 1. *Given  $\mathbb{S} = \{\mathbb{I}_1, \dots, \mathbb{I}_n\}$  and a target  $\mathbb{P}$ , find a minimal  $\mathbb{S}' \subseteq \mathbb{S}$  such that  $\sigma(\mathbb{P}) \subseteq \sigma(\mathbb{S}')$  and  $|\mathbb{S}'|$  is minimized.*

**Problem Solving.** We solve the above problem as follows.

*Step (1) Converting problem on finite sets.* To handle potentially infinite policies and intents, we define a *label* set  $\mathcal{L} = \{\mathbb{P}, \mathbb{I}_1, \dots, \mathbb{I}_n\}$ , encode each element in  $\mathcal{L}$  using MCP, and partition the request space (i.e., the policy and intents) into ECs using Algorithm 1. Recall that  $\mathcal{M}(\mathbb{P})$  and  $\mathcal{M}(\mathbb{I}_i)$  denote the set of ECs of policy  $\mathbb{P}$  and intent  $\mathbb{I}_i$ , respectively. Thus, Problem 1 is reduced to the following finite variant:

687 PROBLEM 2. Given  $\{\mathcal{M}(\mathbb{I}_1), \dots, \mathcal{M}(\mathbb{I}_n)\}$  and  $\mathcal{M}(\mathbb{P})$ , find a minimal subset  $\mathbb{S}$  such that  $\mathcal{M}(\mathbb{P}) \subseteq \mathcal{M}(\mathbb{S})$ .

689 Step (2) Solving problem via ILP Solver. In our datasets, since each policy and intent maps to only  
 690 a small number of ECs, we can solve Problem 2 exactly using ILP solver[24]. Let  $x_i \in \{0, 1\}$  indicate  
 691 whether intent  $\mathbb{I}_i$  is selected. The ILP formulation for Problem 2 is:

$$693 \quad \text{minimize} \quad \sum_{i=1}^n x_i \quad (9)$$

$$696 \quad \text{subject to} \quad \left( \sum_{i: c_j \in \mathcal{M}(\mathbb{I}_i)} x_i \right) \geq 1, \quad \forall c_j \in \mathcal{M}(\mathbb{P}) \quad (10)$$

$$699 \quad x_i \in \{0, 1\}, \quad \forall i = 1, \dots, n \quad (11)$$

## 701 5 Experiment Setup

702 **Implementation.** We implement *AccessRefinery* with  $\sim 5k$  lines of Java code, which includes three  
 703 major components: *Intent Miner*, *Intent Reducer*, and *Multi-theory Constraint Preprocessor (MCP)*.  
 704 For the MCP, we implement two methods to solve the bit-vector constraints, one based on an SAT  
 705 solver, i.e., MiniSAT [15], and one based on a BDD solver, i.e., JavaBDD [47], and refer to them as  
 706 *AccessRefinery*(MiniSAT) and *AccessRefinery*(JavaBDD), respectively.

707 **Baselines.** To compare with *Intent Miner*, we re-implement *Access Analyzer* proposed by AWS,  
 708 with approximately 5k lines of Java code. [AWS Access Analyzer](#) [4] is not open-sourced, we therefore  
 709 reproduce it for a consistent environment setup. To the best of our knowledge, *Access Analyzer* [4] is the only state-of-the-art intent-mining tool. *Access Analyzer* is built on Zelkova [5],  
 710 which relies on an SMT solver. We also reproduced Zelkova and integrated it into *Access Analyzer*.  
 711 To enable a more comprehensive comparison, we replace the underlying SMT solver with CVC5  
 712 and Z3, respectively. The original paper uses Z3 [11] as the SMT solver, and is efficient for the bit-  
 713 vector theory [42]. CVC5 [6], another SMT solver, generally performs better on string theory [36].  
 714 We refer to these two implementations as *Access Analyzer* (Z3) and *Access Analyzer* (CVC5). More-  
 715 over, Quacky [17] uses the ABC solver [3] as its underlying engine, which is primarily designed  
 716 for model counting. The performance of ABC for satisfiability solving is generally slower than that  
 717 of CVC5 and Z3. Therefore, we do not present results for *Access Analyzer* using ABC as its under-  
 718 lying solver. Additionally, AWS provides an online Command Line Interface (CLI) [1] for *Access*  
 719 *Analyzer*, which we use to validate the correctness of our re-implementations.

720 To compare with our *Intent Reducer*, we also implemented two baselines. Both baselines follow  
 721 the *Intent Reducer* algorithm. The only difference is that they use an SMT solver (Z3 or CVC5) to  
 722 partition the ECs of the policy and intents. We refer to them as *Baseline* (Z3) and *Baseline* (CVC5).

723 **Datasets.** We conduct experiments on both real-world and synthetic datasets.

724 *i). Real-world datasets* include 506 policies collected over one year from a customer of a large  
 725 cloud provider (anonymity requirements). Each policy has 1 to 12 statements, each with 5 to 8 keys  
 726 and 5 to 20 distinct values.

727 *ii) Synthetic datasets* 1 for the correctness experiment. To fully test correctness, these datasets  
 728 cover all possible relationships between values:  $v_1 = v_2$  (equal),  $v_1 \subseteq v_2$  or  $v_2 \subseteq v_1$  (overriding),  
 729  $v_1 \cap v_2 \neq \emptyset$ <sup>2</sup> (overlapping), and  $v_1 \cap v_2 = \emptyset$  (disjoint). Specifically, we synthesized *policies* contain-  
 730 ing two statements, each involving two keys, as two keys suffice to test all value and statement re-  
 731 lationships. Since many value combinations are redundant, we selected 12 representative datasets

732  
 733  
 734 <sup>2</sup>For simplicity, this notation indicates intersection, i.e.,  $v_1 \cap v_2 \neq \emptyset$ ,  $v_1/v_2 \neq \emptyset$ , and  $v_2/v_1 \neq \emptyset$ .

(shown in Table 1) that cover all possible relationships. We also include value complements, denoted by  $\bar{v}_4$ , such as NotResource.

Table 1. Categories of Policies for Correctness Testing

Intent Category	Relation Category	Logical Conditions
<b>Allow &amp; Allow</b>	Equal	$v_1 = v_3, v_2 = v_4$
$\text{Allow}_1[k_1=v_1, k_2=v_2]$ ,	Overriding	$v_1 = v_3, v_2 \subseteq v_4$
$\text{Allow}_2[k_1=v_3, k_2=v_4]$	Overlapping1	$v_1 \cap v_3 \neq \emptyset, v_4 \subseteq v_2$
	Overlapping2	$v_1 \cap v_3 \neq \emptyset, v_2 \cap v_4 \neq \emptyset$
	Disjoint	$v_1 \cap v_3 = \emptyset, v_2 \cap v_4 \neq \emptyset$
<b>Allow &amp; Deny</b>	Equal	$v_1 = v_3, v_2 = v_4$
$\text{Allow}_1[k_1=v_1, k_2=v_2]$ ,	Overrrding1	$v_1 = v_3, v_2 \subseteq v_4$
$\text{Deny}_1[k_1=v_3, k_2=v_4]$	Overrrding2	$v_1 = v_3, v_4 \subseteq v_2$
	Disjoint	$v_1 \cap v_3 = \emptyset, v_2 \cap v_4 \neq \emptyset$
<b>Allow &amp; Deny (Comp.)</b>	Overlapping1	$v_1 \cap v_3 \neq \emptyset, v_4 \subseteq v_2$
$\text{Allow}_1[k_1=v_1, k_2=v_2]$ ,	Overlapping2	$v_1 \cap v_3 \neq \emptyset, v_2 \subseteq v_4$
$\text{Deny}_1[k_1=v_3, k_2=\bar{v}_4]$	Disjoint	$v_1 \cap v_3 = \emptyset, v_2 \cap v_4 \neq \emptyset$

iii) **Synthetic datasets** 2 for the scalability experiment. To avoid performance bias, these datasets are designed to (1) cover all possible relationships between values, (2) use keys consistent with those in real-world datasets, and (3) vary the number of statements to assess scalability. Specifically, we generated datasets with 5 and 6 keys, based on our real-world datasets and prior work [16]. In the 5-Key datasets, three keys (Principal, Action, and Resource) have equal values; one key (SourceArn) has pairwise *disjoint* values; and another key (PrincipalArn) has pairwise *overlapping* values. In the 6-Key datasets, we add an additional key (IpAddress) with pairwise *overriding* values. For both settings, we vary the number of Allow statements from 1 to 15.

All experiments run on a Linux server with two 16-core Intel Xeon Gold 6226R CPUs @2.9GHz and 256G memory. Additionally, the timeout for each experiment is set to 1 hour.

## 6 Experiment Results

In the following, we evaluate *AccessRefinery* to answer six research questions:

### 6.1 Is the re-implementation of *Access Analyzer* valid, and is *AccessRefinery* (including MCP, Intent Miner and Intent Reducer) correct?

6.1.1 *Validity of re-implementation of Access Analyzer*. We treat the AWS Access Analyzer via the CLI API as the ground truth. We verified the validity of our reproduced version from two aspects.

**Execution Time.** The trend of execution time for our reproduced version and the commercial version is consistent for both the 5-Key and 6-Key datasets, as shown in Fig. 10. For the 6-Key dataset with 13 to 15 statements, both versions time out ( $>1$  hour). Although the reproduced and commercial versions do not maintain a perfectly linear relationship, this is expected because the underlying SMT solver uses heuristic search algorithms.

**Result Correctness.** The reproduced version produces intents identical to those of the commercial version as shown in the right of Fig. 12.

6.1.2 *Correctness of AccessRefinery*. We validate the correctness of *AccessRefinery* as the following.

**Correctness of MCP.** To ensure that the satisfiability results produced by MCP and SMT solver are equivalent, we conducted a series of basic Boolean operations tests (available in our source code), as follows: (1) Nested expressions:  $(A \wedge B) \vee (\neg A \wedge C)$ ; (2) Implication chains:  $(A \rightarrow B) \wedge (B \rightarrow$




Fig. 10. The total time of mining intents by using re-implementation of *Access Analyzer* and AWS *Access Analyzer* via the CLI API for synthetic datasets.

C); (3) Distributive law:  $A \wedge (B \vee C)$  vs  $(A \wedge B) \vee (A \wedge C)$ ; (4) De Morgan’s laws:  $\neg(A \wedge B)$  vs  $\neg A \vee \neg B$ ; (5) Tautology and contradiction:  $A \vee \neg A$  (tautology) and  $A \wedge \neg A$  (contradiction); (6) Ternary logic / cyclic implications:  $(A \rightarrow B) \wedge (B \rightarrow C) \wedge (C \rightarrow A)$ , verified against  $A \leftrightarrow B \leftrightarrow C$ ; (7) Bi-implication:  $(A \wedge B) \leftrightarrow (C \vee \neg A)$ ; (8) All variables false:  $\neg A \wedge \neg B \wedge \neg C$ ; (9) All variables true:  $A \wedge B \wedge C$ ; (10) Complex combination:  $(A \vee B) \wedge (\neg A \vee C) \wedge (\neg B \vee \neg C)$ . All tests successfully verified the expected satisfiability behaviors, confirming that MCP correctly computes Boolean satisfiability.

**Correctness of Intent Miner.** We compared the intents produced by *AccessRefinery* (without intent reduction), our re-implementation of *Access Analyzer*, and the AWS *Access Analyzer* via the CLI API. On synthetic datasets, all three produce the same set of intents. On real-world datasets, due to privacy constraints, we only compare *AccessRefinery* with our re-implementation of *Access Analyzer*, and they also produce the same set of intents. Fig. 11 shows the number of intents returned by these tools, which further confirms the correctness of MCP.




Fig. 11. The number of intents generated by *AccessRefinery*, our re-implementation of *Access Analyzer*, and AWS *Access Analyzer* via the CLI API for mining intents.

**Correctness of Intent Reducer.** The oracle may use an SMT solver to enumerate all intent sets covering the original policy and select the minimal one. However, since this is slow and even time out (> 24 h), we instead use an SMT solver to check that *Intent Reducer* satisfies (for real-world and synthetic datasets 1 & 2): (1) The reduced intents fully cover the policy; (2) Removing any intent from the reduced intents causes the remaining intents to no longer cover the policy.

**Answer to RQ1:** The re-implementation of *Access Analyzer* shows consistent execution time trends and identical intents compared to the AWS *Access Analyzer* via the CLI API. MCP shows same satisfiability with SMT solver. *Intent Miner* produces the same intents as *Access Analyzer*. *Intent Reducer* produces the minimal intents that still cover the original policy.

## 6.2 Can *AccessRefinery* reduce the number of intents?

Fig. 12 shows that, for real-world datasets, the number of intents can reach up to 80, while for synthetic datasets, it can reach up to 225 intents. These correspond to the original policies containing 5 and 15 statements, respectively. For both real-world and synthetic datasets, *AccessRefinery* reduces the number of intents by up to one order of magnitude. In both types of datasets, the reduction becomes more significant as the number of Allow statements increases.




Fig. 12. The number of intents generated before *reducing intents* and after *reducing intents*. Note that the x-axis for the synthetic datasets represents the number of allowed statements.

**Answer to RQ2:** *AccessRefinery* can reduce the number of intents by up to one order of magnitude for both real and synthetic datasets.

### 6.3 Can *AccessRefinery* speed up intent mining and reduction by using MCP?

For the *intent mining*, Fig. 13 shows that *AccessRefinery* is faster than *Access Analyzer* (Z3) by one to two orders of magnitude, and faster than *Access Analyzer* (CVC5) by one to three orders of magnitude on both the 5-Key and 6-Key datasets. Although *Access Analyzer* (CVC5) is initially faster than *Access Analyzer* (Z3), its performance degrades as the number of allowed statements increases. This is because Z3 handles formulas involving set intersections and bit-vector[42] theory more efficiently on our datasets.




Fig. 13. The total time of mining intents for synthetic datasets.

For the *intent reduction*, Fig. 14 shows that *AccessRefinery* outperforms both *Baseline*(Z3) and *Baseline*(CVC5) by one to four orders of magnitude on the 5-Key and 6-Key datasets. Note that *Baseline*(CVC5) times out with 12 statements on the 5-Key dataset and 8 statements on the 6-Key dataset, while *AccessRefinery* can still reduce intents within ~10s even with 15 statements.





Fig. 14. The total time of reducing intents for synthetic datasets.

**Answer to RQ3:** For the *intent mining*, *AccessRefinery* achieves a speedup of one to two orders of magnitude, and for the *intent reduction*, it achieves a speedup of one to four orders of magnitude.

#### 883 6.4 How does *AccessRefinery* perform on real-world datasets?

884 Since we have previously analyzed the time for individual policies, we now present the cumulative  
 885 time to better evaluate how *AccessRefinery* handles a large number of policies in real-world  
 886 scenarios. Fig. 15 compares the cumulative time for both the *intent mining* and *intent reduction*  
 887 stages. Using Z3, *Access Analyzer* (Z3) and *Baseline* (Z3) require approximately 5983.6s and 2109.1s,  
 888 respectively, to process 506 policies. With CVC5, *Access Analyzer* (CVC5) and *Baseline* (CVC5) still require around 912.3s and  
 889 3401.1s, respectively. In contrast, *AccessRefinery* completes both stages in just 64.2 and 20.2s, re-  
 890 spectively. We clearly observe performance spikes in *Access Analyzer* (Z3) and *Baseline* (Z3) during  
 891 both stages. This is primarily due to the presence of many website-related values, such as patterns  
 892 like “www.example.”, which are expensive to process.  
 893




904 Fig. 15. The cumulative time of the *intent mining* stage and the *intent reduction* stage for real-world datasets.  
 905 The runtime of *AccessRefinery* is not visible in the figure as it is minimal: 64.2s and 20.2s, respectively.  
 906

907 **Answer to RQ4:** For the real-world dataset of 506 policies, *AccessRefinery* achieves substantial  
 908 speedups over the baselines, being up to 15× faster in the *intent mining* and up to 170× faster in  
 909 the *intent reduction*.

#### 912 6.5 Is SAT or BDD better for intent mining and reduction?

913 For *intent mining*, using JavaBDD is 1–6× faster than using MiniSAT<sup>3</sup> (For clarity, the figure is  
 914 omitted.). This improvement mainly comes from solving Formula (4). For *intent reduction*, using  
 915 JavaBDD allows our method two orders of magnitude faster than using MiniSAT, as shown in Fig.  
 916 16. Moreover, when processing 14 statements under the 5-Keys and 6-Keys settings, the MiniSAT-  
 917 based version times out, while the JavaBDD-based version finishes in 100s.



927 Fig. 16. The total time of *reducing intents* for synthetic datasets.

928 <sup>3</sup>Incrementally computing ECs requires clause deletion, which is not efficiently supported by existing approaches. Main-  
 929 taining derivation information for learned clauses [30, 35] or using assumption variables [14, 20] both incur non-negligible  
 930 overhead. Therefore, we do not use an incremental setting in the MiniSAT solver.

**Answer to RQ5:** Our method based on BDD is orders of magnitude faster than that based on SAT for both intent mining and reduction.

## 6.6 How does *AccessRefinery* accelerate single-round solving in multi-round SMT solving compared to SMT solvers?

Due to the limited scalability of SMT solvers, we applied a minor optimization to reduce the number of solving rounds. In the 5-Key dataset, three of the keys share identical value domains; therefore, we assigned specific values to these domains for the root intent in *Intent Miner*, instead of using wildcards (\*). It was not applied to *AccessRefinery*, which already achieves high efficiency without additional optimization.

Table 2. Average Time of Single-Round Boolean Solving (Z3, CVC5, MCP) and Preprocessing Time of MCP on the 5-Key Dataset.

# of allow statements	# of rounds	The time of single-round Boolean solving			The time of MCP preprocessing
		Z3	CVC5	MCP	
3	64	192.1ms	27.9ms	77.5 $\mu$ s	54.5ms
6	196	756.2ms	254.3ms	64.4 $\mu$ s	189.7ms
9	400	1517.2ms	987.0ms	104.1 $\mu$ s	394.3ms
12	676	3122.2ms	4979.3ms	190.2 $\mu$ s	1011.7ms
15	1024	4545.8ms	5543.4ms	274.5 $\mu$ s	1741.5ms

Table 2 shows that for the SMT solver Z3, the average solving time per round ranges from 192.1ms to 4546.8ms, while for CVC5 it ranges from 27.9ms to 5543.4ms. In contrast, *AccessRefinery* employs MCP that extracts theory redundancy in 54.5ms to 1741.5ms. After preprocessing, the average solving time per round is reduced to only 77.5–274.5 $\mu$ s. These results demonstrate that pre-extracting theory redundancy significantly enhances the efficiency of each solving step. Furthermore, Table 2 reveals that the number of solving rounds increases sharply with the number of policy statements, further highlighting the scalability advantage of *AccessRefinery*. Note that although our redundancy extraction is faster than CVC5 in a single round, this does not necessarily mean that we outperform CVC5. The main reason is that real-world regular expressions are generally not very complex. CVC5 may be more efficient when handling very complex string constraints.

**Answer to RQ6:** Through preprocessing SMT constraints, *AccessRefinery* reduces the time of single-round SMT solving from seconds to milliseconds.

## 7 Discussion

In this part, we discuss the limitations and scopes of Multi-Theory Constraint Preprocessor (MCP).

**Limitations.** MCP significantly improves performance over SMT [12] solvers, but it has two main limitations. First, MCP is mainly designed for multi-round SMT solving with predefined values. If a problem only needs to be solved once, MCP may not outperform an SMT solver. Second, the theories supported by MCP are still limited. It only supports a limited set of types, including strings, bit-vectors, ranges, and enumerations. Also, it currently does not support string concatenation, arithmetic, or quantified logic [7]. Extending MCP to support more theories and logics is a valuable direction for future work.

**Scopes.** Even though MCP currently targets specific theories, we tentatively think that MCP can be applied to domains beyond IAM, including: (1) Access Control Lists (ACLs) [50] that filter or

981 redirect packets based on five tuples (source/destination port, source/destination IP, protocol number); and (2) routing policies [39] that filter or modify routes based on a set of attributes (AS path, 982 community tags, local preference, etc.). For example, MCP supports the modeling of ACLs by 983 representing ports as ranges, IPs as bit-vectors, and protocol numbers as enumerations. Moreover, 984 analyzing ACLs for misconfigurations or redundancies also involves multi-round solving [50], and 985 can potentially benefit from using MCP.

## 988 8 Conclusion

989 We propose *AccessRefinery* to fast mine concise intents from complex cloud *IAM policies*. Specifically, 990 *AccessRefinery* uses *Multi-Theory Constraint Preprocessor* to accelerate multi-round SMT solving 991 and applies *Intent Miner* to mine the raw intents and *Intent Reducer* to reduce intents from 992 the raw intents. We evaluate *AccessRefinery* on both real-world and synthetic datasets against 993 the state-of-the-art approach. Experimental results show that *AccessRefinery* achieves  $\sim 10\text{-}100\times$  994 speedup compared to SMT-based methods, while reducing the number of mined intents  $\sim 10\times$ .

## 996 Data Availability

997 The replication package is available at <https://anonymous.4open.science/r/accessrefinery-fse26/>, 998 including the implementation of *AccessRefinery* and the reproduced implementation of *Access Analyzer*, 999 and two synthetic datasets derived from real-world cases. We cannot release the real-world 1000 datasets due to customer permissions, as they contain commercially confidential information.

## 1002 References

- 1004 [1] Amazon Web Services. 2020. AWS IAM Access Analyzer Samples. <https://github.com/aws-samples/aws-iam-access-analyzer-samples>. Accessed: 2025-05-22.
- 1005 [2] Henrik Reif Andersen. 1997. An introduction to binary decision diagrams. *Lecture notes, available online, IT University of Copenhagen* 5 (1997).
- 1007 [3] Abdulbaki Aydin, William Eiers, Lucas Bang, Tegan Brennan, Miroslav Gavrilov, Tevfik Bultan, and Fang Yu. 2018. Parameterized model counting for string and numeric constraints. In *Proceedings of the 2018 26th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering*. 400–410.
- 1009 [4] John Backes, Ulises Berrueco, Tyler Bray, Daniel Brim, Byron Cook, Andrew Gacek, Ranjit Jhala, Kasper Luckow, Sean McLaughlin, Madhav Menon, et al. 2020. Stratified abstraction of access control policies. In *Computer Aided Verification: 32nd International Conference, CAV 2020, Los Angeles, CA, USA, July 21–24, 2020, Proceedings, Part I* 32. Springer, 165–176.
- 1013 [5] John Backes, Pauline Bolignano, Byron Cook, Catherine Dodge, Andrew Gacek, Kasper Luckow, Neha Rungta, Oksana Tkachuk, and Carsten Varming. 2018. Semantic-based automated reasoning for AWS access policies using SMT. In *2018 Formal Methods in Computer Aided Design (FMCAD)*. IEEE, 1–9.
- 1016 [6] Haniel Barbosa, Clark Barrett, Martin Brain, Gereon Kremer, Hanna Lachnitt, Makai Mann, Abdalrhman Mohamed, Mudathir Mohamed, Aina Niemetz, Andres Nötzli, et al. 2022. cvc5: A versatile and industrial-strength SMT solver. In *International Conference on Tools and Algorithms for the Construction and Analysis of Systems*. Springer, 415–442.
- 1019 [7] Mordechai Ben-Ari. 2012. *Mathematical logic for computer science*. Springer Science & Business Media.
- 1022 [8] Malik Bouchet, Byron Cook, Bryant Cutler, Anna Druzkina, Andrew Gacek, Liana Hadarean, Ranjit Jhala, Brad Marshall, Dan Peebles, Neha Rungta, et al. 2020. Block public access: trust safety verification of access control policies. In *Proceedings of the 28th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering*. 281–291.
- 1023 [9] Vasek Chvatal. 1979. A greedy heuristic for the set-covering problem. *Mathematics of operations research* 4, 3 (1979), 233–235.
- 1024 [10] Martin Davis and Hilary Putnam. 1960. A computing procedure for quantification theory. *Journal of the ACM (JACM)* 7, 3 (1960), 201–215.
- 1026 [11] Leonardo De Moura and Nikolaj Bjørner. 2008. Z3: An efficient SMT solver. In *International conference on Tools and Algorithms for the Construction and Analysis of Systems*. Springer, 337–340.
- 1027 [12] Leonardo De Moura and Nikolaj Bjørner. 2011. Satisfiability modulo theories: introduction and applications. *Commun. ACM* 54, 9 (2011), 69–77.

- [13] Loris D'Antoni, Andrew Gacek, Amit Goel, Dejan Jovanovic, Rami Gökhane Kıcı, Dan Peebles, Neha Rungta, Yasmine Sharoda, and Chungha Sung. 2024. Projective model counting for IP addresses in access control policies. In # *PLACEHOLDER\_PARENT\_METADATA\_VALUE#*. TU Wien Academic Press, 208–216.
- [14] N. Een. 2004. An Extensible SAT-solver. In *Proc Sixth International Conference on Theory and Applications of Satisfiability Testing*.
- [15] Niklas Eén and Niklas Sörensson. 2003. An extensible SAT-solver. In *International conference on theory and applications of satisfiability testing*. Springer, 502–518.
- [16] William Eiers, Ganesh Sankaran, and Tevfik Bultan. 2023. Quantitative Policy Repair for Access Control on the Cloud. In *Proceedings of the 32nd ACM SIGSOFT International Symposium on Software Testing and Analysis*. 564–575.
- [17] William Eiers, Ganesh Sankaran, Albert Li, Emily O'Mahony, Benjamin Prince, and Tevfik Bultan. 2022. Quacky: Quantitative Access Control Permissiveness Analyzer. In *Proceedings of the 37th IEEE/ACM International Conference on Automated Software Engineering*. 1–5.
- [18] William Eiers, Ganesh Sankaran, Albert Li, Emily O'Mahony, Benjamin Prince, and Tevfik Bultan. 2022. Quantifying permissiveness of access control policies. In *Proceedings of the 44th International Conference on Software Engineering*. 1805–1817.
- [19] Michael D. Ernst, Jeff H. Perkins, Philip J. Guo, Stephen McCamant, Carlos Pacheco, Matthew S. Tschantz, and Chen Xiao. 2007. The Daikon system for dynamic detection of likely invariants. *Science of Computer Programming* 69, 1 (2007), 35–45. [doi:10.1016/j.scico.2007.01.015](https://doi.org/10.1016/j.scico.2007.01.015) Special issue on Experimental Software and Toolkits.
- [20] Niklas Eén and Niklas Sörensson. 2003. Temporal Induction by Incremental SAT Solving. *Electronic Notes in Theoretical Computer Science* 89, 4 (2003), 543–560.
- [21] Kathi Fisler, Shriram Krishnamurthi, Laurie A. Meyer, and Michael F. Reidy. 2005. Verification and Change-Impact Analysis of Access-Control Policies. In *Proceedings of the 27th International Conference on Software Engineering (ICSE)*. ACM, St. Louis, Missouri, USA, 196–205. [doi:10.1145/1062455.1062481](https://doi.org/10.1145/1062455.1062481)
- [22] Mark Gabel and Zhendong Su. 2008. Symbolic mining of temporal specifications. In *2008 ACM/IEEE 30th International Conference on Software Engineering*, 51–60. [doi:10.1145/1368088.1368096](https://doi.org/10.1145/1368088.1368096)
- [23] Harald Ganzinger, George Hagen, Robert Nieuwenhuis, Albert Oliveras, and Cesare Tinelli. 2004. DPLL (T): Fast decision procedures. In *Computer Aided Verification: 16th International Conference, CAV 2004, Boston, MA, USA, July 13–17, 2004. Proceedings 16*. Springer, 175–188.
- [24] Ralph E Gomory. 2009. Outline of an algorithm for integer solutions to linear programs and an algorithm for the mixed integer problem. In *50 Years of integer programming 1958–2008: From the early years to the State-of-the-Art*. Springer, 77–103.
- [25] Antonios Gouglidis, Anna Kagia, and Vincent C Hu. 2023. Model checking access control policies: A case study using google cloud iam. *arXiv preprint arXiv:2303.16688* (2023).
- [26] Dorit S Hochbaum. 1982. Approximation algorithms for the set covering and vertex cover problems. *SIAM Journal on computing* 11, 3 (1982), 555–556.
- [27] John E Hopcroft, Rajeev Motwani, and Jeffrey D Ullman. 2001. Introduction to automata theory, languages, and computation. *Acm Sigact News* 32, 1 (2001), 60–65.
- [28] Alex Horn, Ali Kheradmand, and Mukul Prasad. 2017. Delta-net: Real-time network verification using atoms. In *14th USENIX Symposium on Networked Systems Design and Implementation (NSDI 17)*. 735–749.
- [29] Richard M Karp. 2009. Reducibility among combinatorial problems. In *50 Years of Integer Programming 1958–2008: from the Early Years to the State-of-the-Art*. Springer, 219–241.
- [30] Joonyoung Kim, J. Whittemore, and K. Sakallah. 2000. On solving stack-based incremental satisfiability problems. In *IEEE International Conference on Computer Design: Vlsi in Computers and Processors*.
- [31] Matija Kuprešanin and Pavle Subotić. 2023. Ambit: Verification of Azure RBAC. In *Proceedings of the 2023 on Cloud Computing Security Workshop*. 31–40.
- [32] Charles Eric Leiserson, Ronald L Rivest, Thomas H Cormen, and Clifford Stein. 1994. *Introduction to algorithms*. Vol. 3. MIT press Cambridge, MA, USA.
- [33] Zechun Li, Peng Zhang, Yichi Zhang, and Hongkun Yang. 2025. {NDD}: A Decision Diagram for Network Verification. In *22nd USENIX Symposium on Networked Systems Design and Implementation (NSDI 25)*. 237–258.
- [34] Xu Liu, Peng Zhang, Hao Li, and Wenbing Sun. 2023. Modular Data Plane Verification for Compositional Networks. *Proceedings of the ACM on Networking* 1, CoNEXT3 (2023), 1–22.
- [35] J. P. Marques-silva and K. A. Sakallah. 1999. GRASP: a search algorithm for propositional satisfiability. *IEEE Trans. Comput.* 48, 5 (1999), 506–521.
- [36] Andres Nötzli, Andrew Reynolds, Haniel Barbosa, Clark Barrett, and Cesare Tinelli. 2022. Even faster conflicts and lazier reductions for string solvers. In *International Conference on Computer Aided Verification*. Springer, 205–226.

- 1079 [37] Port Swigger. 2021. Data leak Down Under: 50,000 gov't employee records found on open S3 bucket.  
1080 https://portswigger.net/daily-swig/data-leak-down-under-50-000-govt-employee-records-found-on-open-s3-  
1081 bucket.
- 1082 [38] Michael O Rabin and Dana Scott. 1959. Finite automata and their decision problems. *IBM journal of research and development* 3, 2 (1959), 114–125.
- 1083 [39] Yakov Rekhter, Susan Hares, and Tony Li. 2006. A Border Gateway Protocol 4 (BGP-4). RFC 4271.
- 1084 [40] SC Media. 2017. Another misconfigured Amazon S3 server leaks data of 50,000 Australian employ-  
1085 ees. https://www.scworld.com/news/another-misconfigured-amazon-s3-server-leaks-data-of-50000-australian-  
1086 employees.
- 1087 [41] Ilia Shevrin and Oded Margalit. 2023. Detecting Multi-Step IAM attacks in AWS environments via model checking.  
In *32nd USENIX Security Symposium (USENIX Security 23)*. 6025–6042.
- 1088 [42] Jiaxin Song, Hongfei Fu, and Charles Zhang. 2024. Equational Bit-Vector Solving via Strong Gr\"obner Bases. *arXiv preprint arXiv:2402.16314* (2024).
- 1089 [43] Robert Roth Stoll. 1979. *Set theory and logic*. Courier Corporation.
- 1090 [44] The Verge. 2021. Microsoft Azure Cloud Vulnerability Is the 'Worst You Can Imagine'.  
https://www.theverge.com/2021/8/27/22644161/microsoft-azure-database-vulnerability-chaosdb.
- 1091 [45] Grigori S Tseitin. 1983. On the complexity of derivation in propositional calculus. *Automation of reasoning: 2: Classical  
1092 papers on computational logic 1967–1970* (1983), 466–483.
- 1093 [46] UpGuard. 2017. Cloud Leak: WSJ Parent Company Dow Jones Exposed Customer Data.  
https://www.upguard.com/breaches/cloud-leak-dow-jones.
- 1094 [47] John Whaley. [n. d.]. JavaBDD: A Java library for Binary Decision Diagrams. <https://javabdd.sourceforge.net/>. Accessed: 2025-05-27.
- 1095 [48] Laurence A Wolsey and George L Nemhauser. 1999. *Integer and combinatorial optimization*. John Wiley & Sons.
- 1096 [49] Hongkun Yang and Simon S Lam. 2015. Real-time verification of network properties using atomic predicates.  
*IEEE/ACM Transactions on Networking* 24, 2 (2015), 887–900.
- 1097 [50] Lihua Yuan, Jianning Mai, Zhendong Su, Hao Chen, and Prasant Mohapatra. 2006. FIREMAN: a toolkit for firewall  
1098 modeling and analysis. *IEEE* (2006).
- 1099 [51] Peng Zhang, Xu Liu, Hongkun Yang, Ning Kang, Zhengchang Gu, and Hao Li. 2020. {APKeep}: Realtime Verification  
1100 for Real Networks. In *17th USENIX Symposium on Networked Systems Design and Implementation (NSDI 20)*. 241–255.
- 1101
- 1102
- 1103
- 1104
- 1105
- 1106
- 1107
- 1108
- 1109
- 1110
- 1111
- 1112
- 1113
- 1114
- 1115
- 1116
- 1117
- 1118
- 1119
- 1120
- 1121
- 1122
- 1123
- 1124
- 1125
- 1126
- 1127